


Telco Customer Churn Prediction

M V Murali Krishna Nagulla

Introduction

Churn prediction is one of the most popular Big Data use cases in business. It consists of detecting customers who are likely to cancel a subscription to a service.

Churn is a problem for telecom companies because it is more expensive to acquire a new customer than to keep your existing one from leaving.

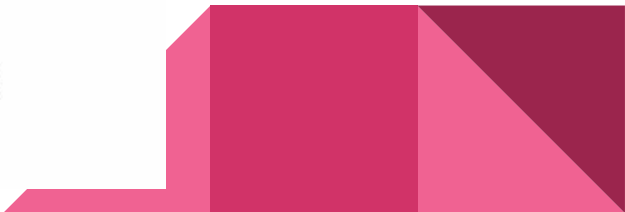


Overview

- **Goal:** Analyze telecommunication company customer data to predict whether or not a customer is likely to leave the platform (churn)
- Data from 7043 customers (21 features):
 - Churn (Yes or No)
 - Customer account information (tenure, contract, payments, etc.)
 - Demographic Information (Partner, Gender, Age, etc.)
 - Add-on services provided by the platform
- Source: [Kaggle.com](https://www.kaggle.com)



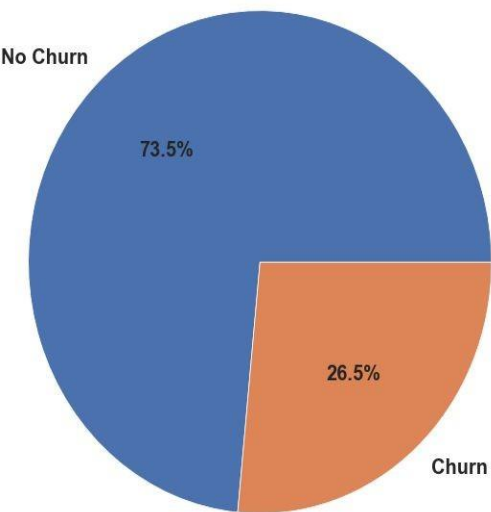
Project Objective

- To predict Customer Churn.
 - Highlighting the main variables/factors influencing Customer Churn.
 - Use various ML algorithms to build prediction models, evaluate the accuracy and performance of these models.
 - Finding out the best model for our business case & providing executive summary.
- 

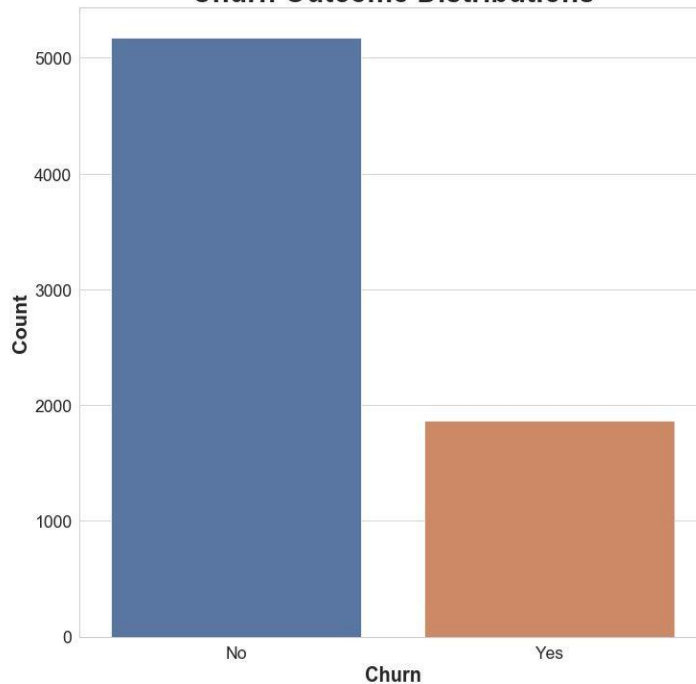
Dataset

- Data
- 20 columns(Removed ID), 7032 entries
- Mostly Categorical Data

Churn Outcome Pie Chart



Churn Outcome Distributions



```
[ ] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 7032 entries, 0 to 7042
```


```
Data columns (total 20 columns):
```

#	Column	Non-Null Count	Dtype
0	gender	7032 non-null	object
1	seniorcitizen	7032 non-null	int64
2	partner	7032 non-null	object
3	dependents	7032 non-null	object
4	tenure	7032 non-null	int64
5	phoneservice	7032 non-null	object
6	multiplelines	7032 non-null	object
7	internetservice	7032 non-null	object
8	onlinesecurity	7032 non-null	object
9	onlinebackup	7032 non-null	object
10	deviceprotection	7032 non-null	object
11	techsupport	7032 non-null	object
12	streamingtv	7032 non-null	object
13	streamingmovies	7032 non-null	object
14	contract	7032 non-null	object
15	paperlessbilling	7032 non-null	object
16	paymentmethod	7032 non-null	object
17	monthlycharges	7032 non-null	float64
18	totalcharges	7032 non-null	float64
19	churn	7032 non-null	object

```
dtypes: float64(2), int64(2), object(16)
```

```
memory usage: 1.1+ MB
```

Feature Correlation to Churn




Most Positively Correlated: Monthly Contracts, No Online Security Add-On


Most Negatively Correlated: Tenure, Two-Year Contracts, No Internet Service

Contract Types


Customer Contract Composition



% Churn - Contract Type




Violin Plot: Monthly Charge - Contract Types




- Significantly more customers will churn when on monthly contracts
- Churn decreases as contract lengths increase
- Customers on monthly contracts are most likely to pay above \$60/bill

Add-On Services




- Customers with online security and/or tech support add-ons will churn the least
- Customers with Streaming services (TV/Movies) will churn the most

Customer Tenure



- Customers are significantly more likely to churn within the first year of tenure on the platform
- As tenure increases, probability of churn decreases

Customer Monthly Charges



- As monthly charges increase, the probability of customer churn increases
- Customers who churn are most likely to have bills exceeding \$60

EDA Conclusions


- Customers with monthly contracts are 20% more likely to churn than with annual contracts
- Customers have the highest probability of churning within the first 20 months on the platform
- Tech-Support & Online Security add-ons play a critical role in preventing churn, while streaming add-ons significantly increase likelihood of churn
- Customers are twice as likely to churn when the monthly charge is greater than \$60



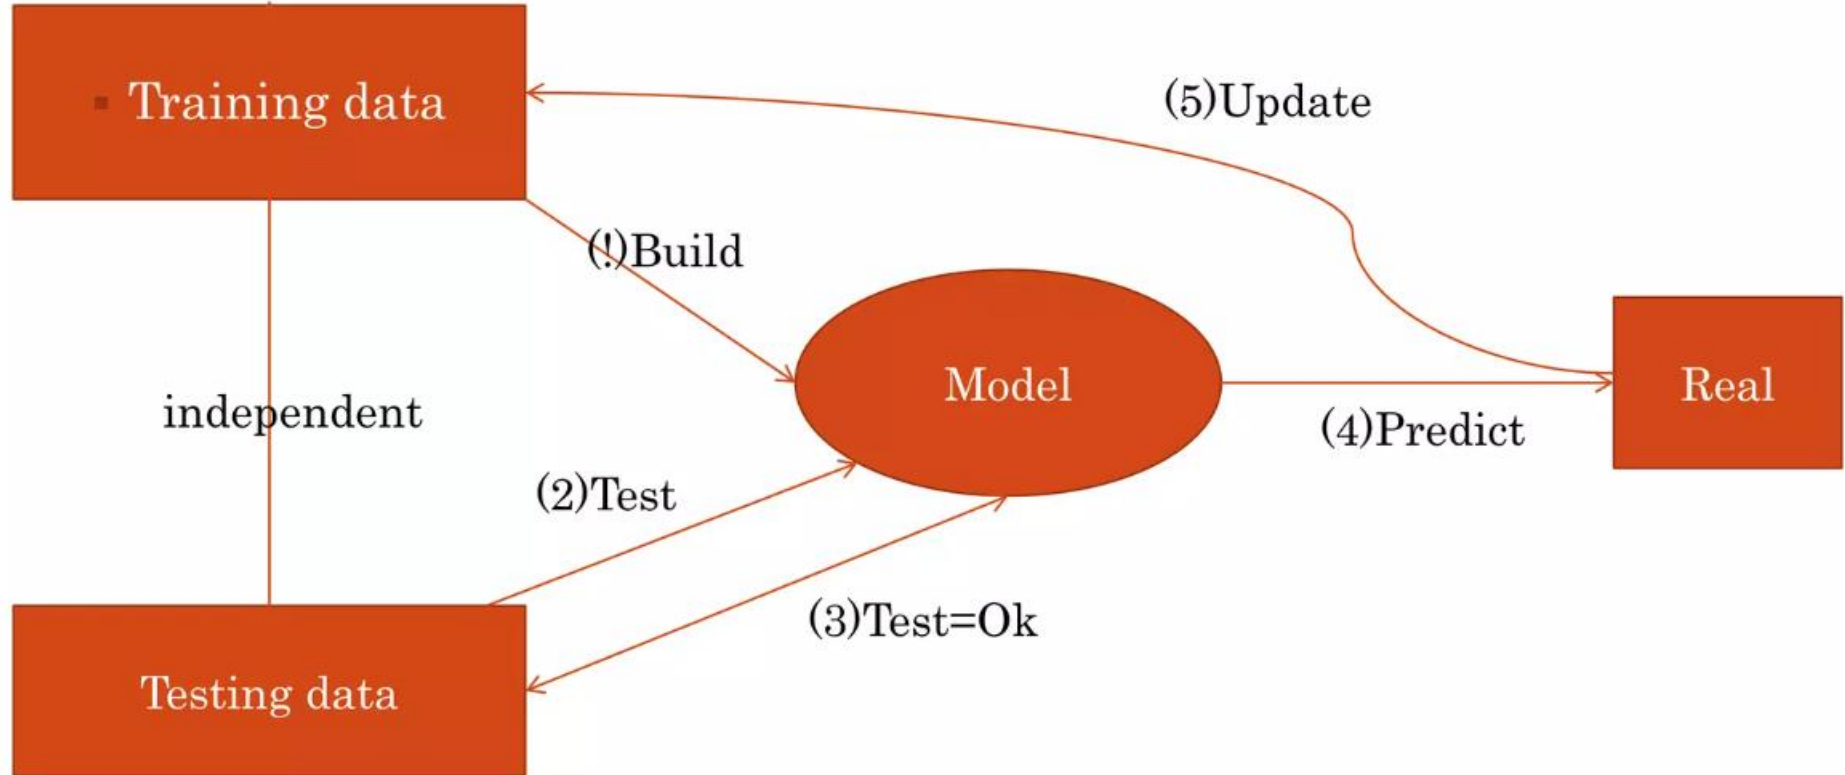
Correlation

- Inference: There is some correlation between 'phone service' and 'multiple lines' since those who don't have a phone service cannot have multiple lines. So, knowing that a particular customer is not subscribed to phone service we can infer that the customer doesn't have multiple lines.
- Similarly, there is also a correlation between 'internet service' and 'online security', 'online backup', 'device protection', 'streaming tv' and 'streaming movies'.





Churn Prediction Model




Data Preprocessing

- One-hot-encoding
- Train Test Split
- Oversampling with SMOTE

One-Hot Encoding

datagy.io

Island	Biscoe	Dream	Torgensen
Biscoe	1	0	0
Torgensen	0	0	1
Dream	0	1	0



SMOTE


SYNTHETIC MINORITY
OVER-SAMPLING TECHNIQUE




Before Smote

	Train Accuracy	Test Accuracy	Overfitting	ROC Area	Precision	Recall	F1-score	Support
CategoricalNB	0.786	0.735	True	0.821049	0.50137	0.783726	0.611529	467
LogisticRegressionCV	0.812	0.759	True	0.828545	0.53429	0.717345	0.612431	467
KNeighborsClassifier	0.798	0.705	True	0.835032	0.469679	0.845824	0.603976	467
DecisionTreeClassifier	0.791	0.742	True	0.78487	0.513514	0.569593	0.540102	467
BaggingClassifier	0.794	0.716	True	0.833156	0.479695	0.809422	0.60239	467
RandomForestClassifier	0.785	0.714	True	0.828506	0.476923	0.796574	0.596632	467
AdaBoostClassifier	0.821	0.759	True	0.838115	0.533435	0.751606	0.624	467
XGBClassifier	0.849	0.763	True	0.815916	0.543403	0.670236	0.600192	467
Linear SVC	0.809	0.762	True	0.826816	0.53871	0.715203	0.614535	467
RBF SVC	0.83	0.77	True	0.817963	0.557798	0.650964	0.600791	467
CatBoostClassifier	0.856	0.768	True	0.829783	0.554128	0.646681	0.596838	467


Smote



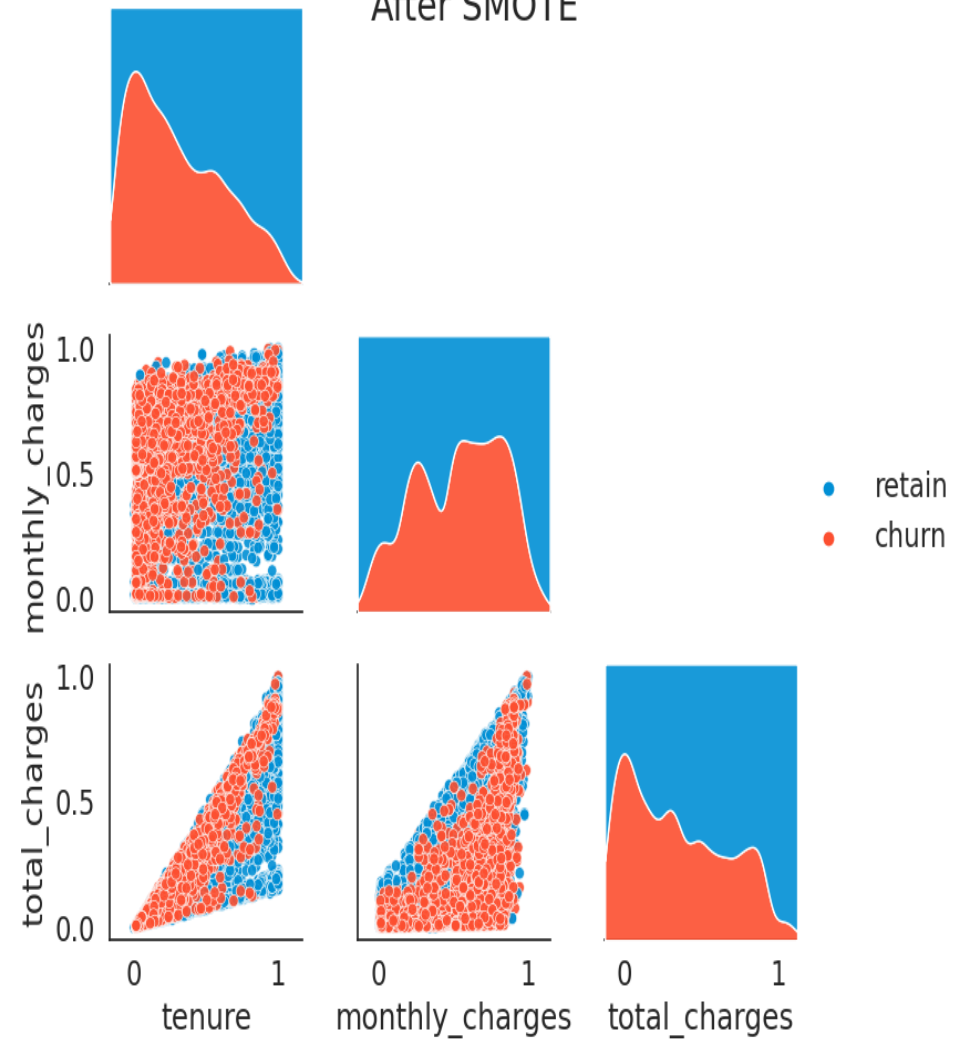
- Majority class samples
- ✚ Minority class samples
- ✚ Randomly selected minority class sample x_i
- ✚ 5 K-nearest neighbors of x_i
- ✚ Randomly selected sample \hat{x}_i from the 5 neighbors
- ✚ Generated synthetic minority instance



Before SMOTE



After SMOTE




Models

	accuracy	macro_avg_precision	macro_avg_recall	macro_avg_f1_score	roc_auc
model					
Logistic Regression	0.746805	0.707463	0.754797	0.714375	0.754797
Ridge Classifier	0.743966	0.706609	0.755140	0.712605	0.755140
KNN	0.696167	0.660272	0.699268	0.661179	0.699268
SVC	0.765736	0.713673	0.747765	0.723961	0.747765
Neural Network	0.758164	0.692219	0.698790	0.695261	0.698790
Decision Tree	0.723616	0.656694	0.671288	0.662195	0.671288
Random Forest	0.773781	0.711190	0.716819	0.713851	0.716819
Gradient Boosting Classifier	0.786559	0.732080	0.758526	0.742016	0.758526
AdaBoost Classifier	0.754851	0.711254	0.755721	0.720050	0.755721
CatBoost Classifier	0.783720	0.723171	0.726430	0.724754	0.726430
Hist Gradient Boosting	0.784193	0.724894	0.734720	0.729374	0.734720
XGBoost	0.778514	0.715981	0.715488	0.715733	0.715488
LightGBM	0.781354	0.720826	0.727665	0.724033	0.727665

Top Performing models

	accuracy	macro_avg_precision	macro_avg_recall	macro_avg_f1_score	roc_auc
model					
Gradient Boosting Classifier	0.786559	0.732080	0.758526	0.742016	0.758526
AdaBoost Classifier	0.754851	0.711254	0.755721	0.720050	0.755721
CatBoost Classifier	0.783720	0.723171	0.726430	0.724754	0.726430
Hist Gradient Boosting	0.784193	0.724894	0.734720	0.729374	0.734720
XGBoost	0.778514	0.715981	0.715488	0.715733	0.715488
LightGBM	0.781354	0.720826	0.727665	0.724033	0.727665

Feature Selection



Comparison


I did model performance comparison before and after feature selection. I'll just take the average of each metrics.

	accuracy	macro_avg_precision	macro_avg_recall	macro_avg_f1_score	roc_auc
original	0.778198	0.721368	0.736425	0.725993	0.736425
filter method	0.772598	0.718596	0.745133	0.726751	0.745133
wrapper method	0.774333	0.719186	0.742710	0.726742	0.742710
embedded method	0.766130	0.712105	0.739781	0.720457	0.739781

Hyperparameter Tuning

	accuracy	precision	recall	f1_score	roc_auc	accuracy	precision	recall	f1_score	roc_auc
model										
Gradient Boosting Classifier	0.786559	0.581602	0.698752	0.634818	0.758526	0.782773	0.570055	0.739750	0.643910	0.769038
AdaBoost Classifier	0.754851	0.526642	0.757576	0.621345	0.755721	0.764789	0.540404	0.762923	0.632668	0.764194
CatBoost Classifier	0.783720	0.590592	0.604278	0.597357	0.726430	0.770469	0.550398	0.739750	0.631179	0.760661
Hist Gradient Boosting	0.784193	0.587354	0.629234	0.607573	0.734720	0.764316	0.539130	0.773619	0.635432	0.767286
XGBoost	0.778514	0.583184	0.581105	0.582143	0.715488	0.777094	0.563025	0.716578	0.630588	0.757773
LightGBM	0.781354	0.584041	0.613191	0.598261	0.727665	0.760530	0.533414	0.782531	0.634393	0.767554

- After tuning, the accuracy score is mostly decreased.
- But, the recall score has increased dramatically. Therefore, I will use the tuned model for model selection.



Model Selection

- F-beta score to calculate the harmonic mean of accuracy and recall.
- Here I use beta=1, that means the accuracy and recall are considered as equally important. If you more care about recall, you can change β to be higher than 1, and vice versa.
- LightGBM

$$F_{\beta} = (1 + \beta^2) \frac{\text{accuracy} * \text{recall}}{\beta * \text{accuracy} + \text{recall}}$$

	accuracy	recall	fbeta
model			
Gradient Boosting Classifier	0.774255	0.762923	0.768547
AdaBoost Classifier	0.759110	0.784314	0.771506
CatBoost Classifier	0.761477	0.768271	0.764859
Hist Gradient Boosting	0.755797	0.782531	0.768932
XGBoost	0.759110	0.748663	0.753850
LightGBM	0.761477	0.786096	0.773591

Conclusion

Final Model

LightGBM with feature selection using filter method

We should pay more attention to customers who meet the criteria below:

Contract: Month-to-month

Tenure: Short tenure

Internet service: Fiber optic

Payment method: Electronic check

Please, evaluate our service!

Especially for internet service (fiber optic) and payment method (electronic check)

Can we give more benefit to a new customer?

Because the new customer has a high probability to churn

